

# Data-X Fall 2018: Homework 06

## Machine Learning

**Authors:** Sana Iqbal (Part 1, 2, 3)

In this homework, you will do some exercises with prediction.

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: # machine learning libraries
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
#import xgboost as xgb
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/ensemble/weight_boosting.p
y:29: DeprecationWarning: numpy.core.umath_tests is an internal NumPy mod
ule and should not be imported. It will be removed in a future NumPy rele
ase.
```

```
from numpy.core.umath_tests import inner1d
```

## Part 1

\_\_ 1. Read **diabetesdata.csv** file into a pandas dataframe. About the data: \_\_

1. **TimesPregnant:** Number of times pregnant
2. **glucoseLevel:** Plasma glucose concentration a 2 hours in an oral glucose tolerance test
3. **BP:** Diastolic blood pressure (mm Hg)
4. **insulin:** 2-Hour serum insulin (mu U/ml)
5. **BMI:** Body mass index (weight in kg/(height in m)<sup>2</sup>)
6. **pedigree:** Diabetes pedigree function
7. **Age:** Age (years)
8. **IsDiabetic:** 0 if not diabetic or 1 if diabetic)

```
In [3]: #Read data & print it
data = pd.read_csv('/Users/chelseayang/Downloads/diabetesdata.csv')
print(data)
data.columns
```

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	\
0	6	148.0	72	0	33.6	0.627	50.0	
1	1	NaN	66	0	26.6	0.351	31.0	
2	8	183.0	64	0	23.3	0.672	NaN	
3	1	NaN	66	94	28.1	0.167	21.0	
4	0	137.0	40	168	43.1	2.288	33.0	
5	5	116.0	74	0	25.6	0.201	30.0	
6	3	78.0	50	88	31.0	0.248	26.0	
7	10	115.0	0	0	35.3	0.134	29.0	
8	2	197.0	70	543	30.5	0.158	53.0	
9	8	NaN	96	0	0.0	0.232	54.0	
10	4	110.0	92	0	37.6	0.191	NaN	
11	10	168.0	74	0	38.0	0.537	34.0	
12	10	139.0	80	0	27.1	1.441	57.0	
13	1	NaN	60	846	30.1	0.398	59.0	
14	5	166.0	72	175	25.8	0.587	51.0	
15	7	100.0	0	0	30.0	0.484	32.0	
16	0	NaN	84	230	45.8	0.551	31.0	
17	7	107.0	74	0	29.6	0.254	31.0	
18	1	100.0	66	0	26.0	0.167	21.0	

## 2. Calculate the percentage of NaN values in each column.

```
In [4]: NullsPerColumn = pd.DataFrame(np.random.randint(low=0, high=1, size=(1, 8)),c

for i in range(0,7):
    perc=round(data.iloc[:,i].isnull().sum()/768,4)*100
    perc=str(perc)+str('%')
    NullsPerColumn.iloc[[0],i]=perc

print(NullsPerColumn)
```

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabet
ic								
0	0.0%	4.43%	0.0%	0.0%	0.0%	0.0%	4.3%	
0								

```
In [5]: ###RUN THIS CELL BUT DO NOT ALTER IT
#assert all(NullsPerColumn.columns == ['Percentage Null'])
#assert NullsPerColumn['Percentage Null'][-2] == 0.04296875
```

## 3. Calculate the TOTAL percent of ROWS with NaN values in the dataframe (make sure values are floats).

```
In [6]: null_val=data.isnull().sum(axis=1)

count_row=0

for i in range(0,768):
    if null_val[i]>=1:
        count_row=count_row+1

percent_null=count_row/768*1.0

percentNull = round(percent_null,4)*100

PercentNull=str(percentNull)+str('%')

print(PercentNull)
```

8.33%

#### 4. Split data into train\_df and test\_df with 15% test split.

```
In [7]: #split values
from sklearn.model_selection import train_test_split

train_df, test_df = train_test_split(data,test_size=0.15,random_state=15)
```

```
In [8]: ###RUN THIS CELL BUT DO NOT ALTER IT
np.testing.assert_almost_equal(float(len(train_df))/float(len(data)), 0.8489)
np.testing.assert_almost_equal(float(len(test_df))/float(len(data)), 0.15104)
```

#### 5. Replace the Nan values in train\_df and test\_df with the mean of EACH feature.

```

In [9]: guess_val_train = []
        guess_val_test=[]

        for i in range(0,8):
            guess_train_i=train_df.iloc[:,i].mean()
            guess_val_train.append(guess_train_i)

        # replace the Nan values in train_df with the mean of EACH feature
        for i in range(0,8):
            train_df.iloc[:,i]=train_df.iloc[:,i].fillna(guess_val_train[i])

        # Replace the nan values in test_df
        for i in range(0,8):
            guess_test_i=test_df.iloc[:,i].mean()
            guess_val_test.append(guess_test_i)

        # replace the Nan values in test_df with the mean of EACH feature
        for i in range(0,8):
            test_df.iloc[:,i]=test_df.iloc[:,i].fillna(guess_val_test[i])

        print(train_df)
        print(test_df)

```

/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py:543: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)  
self.obj[item] = s

```

In [10]: ###RUN THIS CELL BUT DO NOT ALTER IT

        #assert sum(train_df.isnull().sum()) == 0
        #assert sum(test_df.isnull().sum()) == 0

```

**6. Split train\_df & test\_df into x\_train, y\_train and x\_test, y\_test .**  
**y\_train and y\_test should only have the column we are trying to predict, IsDiabetic .**

```
In [11]: X_train = train_df.drop('IsDiabetic',axis=1)
Y_train = train_df['IsDiabetic']
X_test  = test_df.drop('IsDiabetic',axis=1)
Y_test  = test_df['IsDiabetic']
```

```
In [12]: ###RUN THIS CELL BUT DO NOT ALTER IT
assert [X_train.shape, Y_train.shape, X_test.shape,Y_test.shape] == [(652,
```

**7. Use this dataset to train perceptron, logistic regression and random forest models using 15% test split. Report training and test accuracies.**

```
In [13]: # Logistic Regression

logreg = LogisticRegression()
logreg.fit(X_train,Y_train)
logreg_train_acc =logreg.score(X_train,Y_train)
logreg_test_acc = logreg.score(X_test, Y_test)
print ('logreg training acuracy= ',logreg_train_acc)
print('logreg test accuracy= ',logreg_test_acc)

logreg training acuracy=  0.7822085889570553
logreg test accuracy=  0.7413793103448276
```

```
In [14]: # Perceptron

perceptron = Perceptron()
perceptron.fit(X_train,Y_train)
perceptron_train_acc = perceptron.score(X_train,Y_train)
perceptron_test_acc = perceptron.score(X_test, Y_test)
print ('perceptron training acuracy= ',perceptron_train_acc)
print('perceptron test accuracy= ',perceptron_test_acc)
```

```
perceptron training acuracy=  0.651840490797546
perceptron test accuracy=  0.6982758620689655
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/stochastic_gradient.py:128: FutureWarning: max_iter and tol parameters have been added in <class 'sklearn.linear_model.perceptron.Perceptron'> in 0.19. If both are left unset, they default to max_iter=5 and tol=None. If tol is not None, max_iter defaults to max_iter=1000. From 0.21, default max_iter will be 1000, and default tol will be 1e-3.
```

```
"and default tol will be 1e-3." % type(self), FutureWarning)
```

```
In [15]: # Adaboost
adaboost = AdaBoostClassifier(n_estimators=100)
adaboost.fit(X_train,Y_train)
adaboost_train_acc = adaboost.score(X_train, Y_train)
adaboost_test_acc = adaboost.score(X_test, Y_test)
print('adaboost training accuracy= ',adaboost_train_acc)
print('adaboost test accuracy= ',adaboost_test_acc)
```

```
adaboost training accuracy= 0.8404907975460123
adaboost test accuracy= 0.75
```

```
In [16]: # Random Forest

random_forest = RandomForestClassifier(n_estimators=500)
random_forest.fit(X_train,Y_train)
random_forest_train_acc = random_forest.score(X_train,Y_train)
random_forest_test_acc = random_forest.score(X_test, Y_test)
print('random_forest training accuracy= ',random_forest_train_acc)
print('random_forest test accuracy= ',random_forest_test_acc)
```

```
random_forest training accuracy= 1.0
random_forest test accuracy= 0.75
```

**8. Is mean imputation is the best type of imputation to use? Why or why not? What are some other ways to impute the data?**

Your answer here

```
In [17]: #No, because sometimes the mean of this column is not representative of the
#and if there are too many outliers, the mean will become very bad to
#impute nan value. All in all, firstly, we have to understand what the nan
#and why they are nan, and then decide to use what kind of method.
```

## Part 2

**1.Add columns `BMIband` & `Pedigree\_band` to `Data` by cutting `BMI` & `Pedigree` into 3 intervals. PRINT the first 5 rows of `\_\_data` .**

```
In [18]: # YOUR CODE HERE
#raise NotImplementedError()

data['BMI_band']=pd.cut(data['BMI'],3)

data['Pedigree_band']=pd.cut(data['Pedigree'],3)

data.head()
```

Out[18]:

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic	BMI_band	Pedigree_
0	6	148.0	72	0	33.6	0.627	50.0	1	(22.367, 44.733]	(0.0757, 0.859]
1	1	NaN	66	0	26.6	0.351	31.0	0	(22.367, 44.733]	(0.0757, 0.859]
2	8	183.0	64	0	23.3	0.672	NaN	1	(22.367, 44.733]	(0.0757, 0.859]
3	1	NaN	66	94	28.1	0.167	21.0	0	(22.367, 44.733]	(0.0757, 0.859]
4	0	137.0	40	168	43.1	2.288	33.0	1	(22.367, 44.733]	(1.639, 2.42]

1a. Print the category intervals for `'BMIband'` & `'Pedigreeband'`.

```
In [19]: print('BMI_Band_Interval: ' + str(data['BMI_band'].unique()))

print('Pedigree_Band_Interval: ' + str(data['Pedigree_band'].unique()))

BMI_Band_Interval: [(22.367, 44.733], (-0.0671, 22.367], (44.733, 67.1]]
Categories (3, interval[float64]): [(-0.0671, 22.367] < (22.367, 44.733] < (44.733, 67.1]]
Pedigree_Band_Interval: [(0.0757, 0.859], (1.639, 2.42], (0.859, 1.639]]
Categories (3, interval[float64]): [(0.0757, 0.859] < (0.859, 1.639] < (1.639, 2.42]]
```

2. Group \_\_ data by `Pedigree_band` & determine ratio of diabetic in each band.\_\_

```
In [20]: # YOUR CODE HERE
#raise NotImplementedError()

diabetic_sum = data.groupby('Pedigree_band')['IsDiabetic'].sum()

count=data.groupby('Pedigree_band')['IsDiabetic'].count()

pedigree_DiabeticRatio=diabetic_sum/count*1.0

print(pedigree_DiabeticRatio)
```

```
Pedigree_band
(0.0757, 0.859]    0.327007
(0.859, 1.639]    0.540541
(1.639, 2.42]     0.444444
Name: IsDiabetic, dtype: float64
```

**2a. Group \_\_ data by BMI\_band & determine ratio of diabetic in each band.\_\_**

```
In [21]: # YOUR CODE HERE
#raise NotImplementedError()

diabetic_sum2 = data.groupby('BMI_band')['IsDiabetic'].sum()

count2=data.groupby('BMI_band')['IsDiabetic'].count()

BMI_DiabeticRatio=diabetic_sum2/count2*1.0

print(BMI_DiabeticRatio)
```

```
BMI_band
(-0.0671, 22.367]    0.039216
(22.367, 44.733]    0.358297
(44.733, 67.1]      0.611111
Name: IsDiabetic, dtype: float64
```

```
In [22]: ###RUN THIS CELL BUT DO NOT ALTER IT
#assert BMI_DiabeticRatio['IsDiabetic'][1] == 0.35829662261380324
#assert pedigree_DiabeticRatio['IsDiabetic'][1] == 0.5405405405405406
```

**3. Convert these features - 'BP','insulin','BMI' and 'Pedigree' into categorical values by mapping different bands of values of these features to integers 0,1,2.**

HINT: USE pd.cut with bin=3 to create 3 bins



```
In [23]: # YOUR CODE HERE
#raise NotImplementedError()

data['BP'] = pd.cut(data['BP'], 3, labels=[0, 1, 2])

data['insulin'] = pd.cut(data['insulin'], 3, labels=[0, 1, 2])

data['BMI'] = pd.cut(data['BMI'], 3, labels=[0, 1, 2])

data['Pedigree'] = pd.cut(data['Pedigree'], 3, labels=[0, 1, 2])

data.head()
```

Out[23]:

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic	BMI_band	Pedigree_
0	6	148.0	1	0	1	0	50.0	1	(22.367, 44.733]	(0.0757, 0
1	1	NaN	1	0	1	0	31.0	0	(22.367, 44.733]	(0.0757, 0
2	8	183.0	1	0	1	0	NaN	1	(22.367, 44.733]	(0.0757, 0
3	1	NaN	1	0	1	0	21.0	0	(22.367, 44.733]	(0.0757, 0
4	0	137.0	0	0	1	2	33.0	1	(22.367, 44.733]	(1.639,

```
In [24]: ###RUN THIS CELL BUT DO NOT ALTER IT
#assert sum(data['insulin'])==49
#assert sum(data['BMI'])==753
#assert sum(data['Pedigree'])==92
```

4. Now consider the original dataset again, instead of generalizing the NAN values with the mean of the feature we will try assigning values to NANs based on some hypothesis. For example for age we assume that the relation between BMI and BP of people is a reflection of the age group. We can have 9 types of BMI and BP relations and our aim is to find the median age of each of that group:

Your Age guess matrix will look like this:

BMI	0	1	2
BP			
0	a00	a01	a02
1	a10	a11	a12
2	a20	a21	a22

Create a guess\_matrix for NaN values of 'Age' ( using 'BMI' and 'BP') and 'glucoseLevel' (using 'BP' and 'Pedigree') for the given dataset and assign values accordingly to the NaNs in 'Age' or 'glucoseLevel' .

Refer to how we guessed age in the titanic notebook in the class.

```
In [25]: # YOUR CODE HERE
#raise NotImplementedError()

age_guess=np.zeros((3,3),dtype=int)

for i in range(0,3):
    for j in range(0,3):
        guess_df_age=data[(data['BMI']==i) & (data['BP']==j)][ 'Age' ].dropna()
        guess_age=guess_df_age.median()
        age_guess[i,j]=int(guess_age)

for i in range(0,3):
    for j in range(0,3):
        data.loc[(data.Age.isnull()) & (data['BMI']==i) & (data['BP']==j),

data.Age=data[ 'Age' ].astype(int)
data.head()
```

Out[25]:

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic	BMI_band	Pedigree_I
0	6	148.0	1	0	1	0	50	1	(22.367, 44.733]	(0.0757, 0
1	1	NaN	1	0	1	0	31	0	(22.367, 44.733]	(0.0757, 0
2	8	183.0	1	0	1	0	29	1	(22.367, 44.733]	(0.0757, 0
3	1	NaN	1	0	1	0	21	0	(22.367, 44.733]	(0.0757, 0
4	0	137.0	0	0	1	2	33	1	(22.367, 44.733]	(1.639,

```

In [26]: glucose_guess=np.zeros((3,3),dtype=int)

for i in range(0,3):
    for j in range(0,3):
        guess_df_glucose=data[(data['Pedigree']==i) & (data['BP']==j)][['glucoseLevel']]
        guess_glucose=guess_df_glucose.median()
        glucose_guess[i,j]=int(guess_glucose)

for i in range(0,3):
    for j in range(0,3):
        data.loc[(data['glucoseLevel'].isnull()) & (data['Pedigree']==i) & (data['BP']==j)]['glucoseLevel']=glucose_guess[i,j]

data['glucoseLevel']=data['glucoseLevel'].astype(int)

data.head()

```

Out[26]:

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic	BMI_band	Pedigree_I
0	6	148	1	0	1	0	50	1	(22.367, 44.733]	(0.0757, 0
1	1	112	1	0	1	0	31	0	(22.367, 44.733]	(0.0757, 0
2	8	183	1	0	1	0	29	1	(22.367, 44.733]	(0.0757, 0
3	1	112	1	0	1	0	21	0	(22.367, 44.733]	(0.0757, 0
4	0	137	0	0	1	2	33	1	(22.367, 44.733]	(1.639,

5. Now, convert 'glucoseLevel' and 'Age' features also to categorical variables of 4 categories each. PRINT the head of \_\_ data \_\_

```
In [27]: # YOUR CODE HERE
#raise NotImplementedError()

data['glucoseLevel']=pd.cut(data['glucoseLevel'],4,labels=[0,1,2,3])

data['Age']=pd.cut(data['Age'],4,labels=[0,1,2,3])

data.head()
```

Out[27]:

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic	BMI_band	Pedigree_I
0	6	2	1	0	1	0	1	1	(22.367, 44.733]	(0.0757, 0
1	1	2	1	0	1	0	0	0	(22.367, 44.733]	(0.0757, 0
2	8	3	1	0	1	0	0	1	(22.367, 44.733]	(0.0757, 0
3	1	2	1	0	1	0	0	0	(22.367, 44.733]	(0.0757, 0
4	0	2	0	0	1	2	0	1	(22.367, 44.733]	(1.639,

**6. Use this dataset (with all features in categorical form) to train perceptron, logistic regression and random forest models using 15% test split. Report training and test accuracies.**

```
In [37]: train_df, test_df = train_df, test_df = train_test_split(data, test_size=0.1
X_train = train_df.iloc[:,0:7]
Y_train = train_df.iloc[:,7]
X_test = test_df.iloc[:,0:7]
Y_test = test_df.iloc[:,7]

X_train.shape, Y_train.shape, X_test.shape
```

Out[37]: ((652, 7), (652,), (116, 7))

```
In [38]: # Logistic Regression
logreg = LogisticRegression()
logreg.fit(X_train, Y_train)
logreg_train_acc = logreg.score(X_train, Y_train)
logreg_test_acc = logreg.score(X_test, Y_test)
print('logreg training accuracy= ', logreg_train_acc)
print('logreg test accuracy= ', logreg_test_acc)

logreg training accuracy= 0.7254601226993865
logreg test accuracy= 0.8103448275862069
```

In [39]: *# Perceptron*

```
perceptron = Perceptron()
perceptron.fit(X_train,Y_train)
perceptron_train_acc = perceptron.score(X_train,Y_train)
perceptron_test_acc = perceptron.score(X_test, Y_test)
print ('perceptron training acuracy= ',perceptron_train_acc)
print('perceptron test accuracy= ',perceptron_test_acc)
```

```
perceptron training acuracy=  0.691717791411043
perceptron test accuracy=  0.6724137931034483
```

/anaconda3/lib/python3.7/site-packages/sklearn/linear\_model/stochastic\_gradient.py:128: FutureWarning: max\_iter and tol parameters have been added in <class 'sklearn.linear\_model.perceptron.Perceptron'> in 0.19. If both are left unset, they default to max\_iter=5 and tol=None. If tol is not None, max\_iter defaults to max\_iter=1000. From 0.21, default max\_iter will be 1000, and default tol will be 1e-3.

"and default tol will be 1e-3." % type(self), FutureWarning)

In [40]: *# Random Forest*

```
random_forest = RandomForestClassifier(n_estimators=500)
random_forest.fit(X_train,Y_train)
random_forest_train_acc = random_forest.score(X_train,Y_train)
random_forest_test_acc = random_forest.score(X_test, Y_test)
print('random_forest training acuracy= ',random_forest_train_acc)
print('random_forest test accuracy= ',random_forest_test_acc)
```

```
random_forest training acuracy=  0.8650306748466258
random_forest test accuracy=  0.6982758620689655
```